

Lab Task 4

Name: Tazmeen Afroz
Roll No: 22p-9252
Section: BAI-5A

Basic Insights

```
In [23]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn.preprocessing

df = pd.read_csv('auto-mpg.csv')
df.head()
```

Out[23]:

	mpg	cylinders	displacement	horsepower	weight	acceleration	model year	origin	car name
0	18.0	8	307.0	130	3504	12.0	70	1	chevr chev ma
1	15.0	8	350.0	165	3693	11.5	70	1	b sky
2	18.0	8	318.0	150	3436	11.0	70	1	plymc sate
3	16.0	8	304.0	150	3433	12.0	70	1	rebe
4	17.0	8	302.0	140	3449	10.5	70	1	l to

Column Name	Description	Data Type	Range/Units	Notes
mpg	Miles per gallon	Continuous	9.0 - 46.6	Target variable; measure of fuel efficiency
cylinders	Number of cylinders in the engine	Discrete	3 - 8	Categorical, but represented as integers
displacement	Engine displacement	Continuous	68 - 455	Measured in cubic inches
horsepower	Engine horsepower	Continuous	46 - 230	Contains some missing values (denoted as '?')
weight	Vehicle weight	Continuous	1613 - 5140	Measured in pounds

Column Name	Description	Data Type	Range/Units	Notes
acceleration	Time to accelerate from 0 to 60 mph	Continuous	8.0 - 24.8	Measured in seconds
model year	Model year of the vehicle	Discrete	70 - 82	Represents model years 1970 to 1982
origin	Origin of the vehicle	Categorical	1, 2, 3	1 = USA, 2 = Europe, 3 = Japan
car name	Name of the vehicle	String	N/A	Unique for each instance

Notes:

1. The dataset contains **398 instances** (rows) and **9 attributes** (columns).
2. **'mpg'** is the target variable, typically used for prediction tasks in machine learning models.
3. **'horsepower'** is the only column with missing values, which may require handling in data preprocessing.
4. **'origin'** is encoded as numbers but represents categorical data and might need to be treated as such in analyses.
5. **'car name'** provides additional context but is typically not used as a feature in predictive modeling.
6. The dataset spans car models from **1970 to 1982**, capturing a significant period in automotive history.
7. There's a mix of continuous and categorical variables, which may require different preprocessing techniques.
8. Some variables (like **'weight'** and **'horsepower'**) may have strong correlations with the target variable **'mpg'**.

```
In [24]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   mpg              398 non-null   float64
1   cylinders        398 non-null   int64
2   displacement     398 non-null   float64
3   horsepower       398 non-null   object
4   weight           398 non-null   int64
5   acceleration     398 non-null   float64
6   model year      398 non-null   int64
7   origin           398 non-null   int64
8   car name        398 non-null   object
dtypes: float64(3), int64(4), object(2)
memory usage: 28.1+ KB
```

```
In [25]: df.describe(include="all").T.round(2)
```

Out[25]:

	count	unique	top	freq	mean	std	min	25%	50%
mpg	398.0	NaN	NaN	NaN	23.514573	7.815984	9.0	17.5	21.0
cylinders	398.0	NaN	NaN	NaN	5.454774	1.701004	3.0	4.0	4.0
displacement	398.0	NaN	NaN	NaN	193.425879	104.269838	68.0	104.25	146.0
horsepower	398	94	150	22	NaN	NaN	NaN	NaN	NaN
weight	398.0	NaN	NaN	NaN	2970.424623	846.841774	1613.0	2223.75	2800.0
acceleration	398.0	NaN	NaN	NaN	15.56809	2.757689	8.0	13.825	16.0
model year	398.0	NaN	NaN	NaN	76.01005	3.697627	70.0	73.0	76.0
origin	398.0	NaN	NaN	NaN	1.572864	0.802055	1.0	1.0	1.0
car name	398	305	ford pinto	6	NaN	NaN	NaN	NaN	NaN

```
In [26]: numerical_vars = df.select_dtypes(include=['int64', 'float64']).columns.tolist()
categorical_vars = df.select_dtypes(include=['object']).columns.tolist()
print('Numerical variables:', numerical_vars)
print('Categorical variables:', categorical_vars)
```

Numerical variables: ['mpg', 'cylinders', 'displacement', 'weight', 'acceleration', 'model year', 'origin']
Categorical variables: ['horsepower', 'car name']

```
In [27]: non_numeric = df['horsepower'][pd.to_numeric(df['horsepower'], errors='coerce').isna()]
print(non_numeric.unique())
```

['?']

```
In [28]: # Replace '?' with NaN
df['horsepower'] = df['horsepower'].replace('?', np.nan)

df['horsepower'] = pd.to_numeric(df['horsepower'])
```

```
In [29]: categorical_count = df.select_dtypes(include='object').shape[1]
numerical_count = df.select_dtypes(exclude='object').shape[1]

print(f"Number of categorical variables: {categorical_count}")
print(f"Number of numerical variables: {numerical_count}")
```

Number of categorical variables: 1
Number of numerical variables: 8

Visualization

```
In [30]: import seaborn as sns
import matplotlib.pyplot as plt

sns.set_style("whitegrid")

plt.figure(figsize=(16, 12))
```

```
# Subplot 1 - MPG (Miles per Gallon)
plt.subplot(3, 3, 1)
sns.histplot(df['mpg'], bins=30, kde=True, color='teal')
plt.title('Miles per Gallon (MPG)', fontsize=14)
plt.xlabel('MPG', fontsize=12)
plt.ylabel('Frequency', fontsize=12)

# Subplot 2 - Displacement
plt.subplot(3, 3, 2)
sns.histplot(df['displacement'], bins=30, kde=True, color='coral')
plt.title('Engine Displacement', fontsize=14)
plt.xlabel('Displacement (cubic inches)', fontsize=12)
plt.ylabel('Frequency', fontsize=12)

# Subplot 3 - Horsepower
plt.subplot(3, 3, 3)
sns.histplot(df['horsepower'], bins=30, kde=True, color='dodgerblue')
plt.title('Horsepower', fontsize=14)
plt.xlabel('Horsepower', fontsize=12)
plt.ylabel('Frequency', fontsize=12)

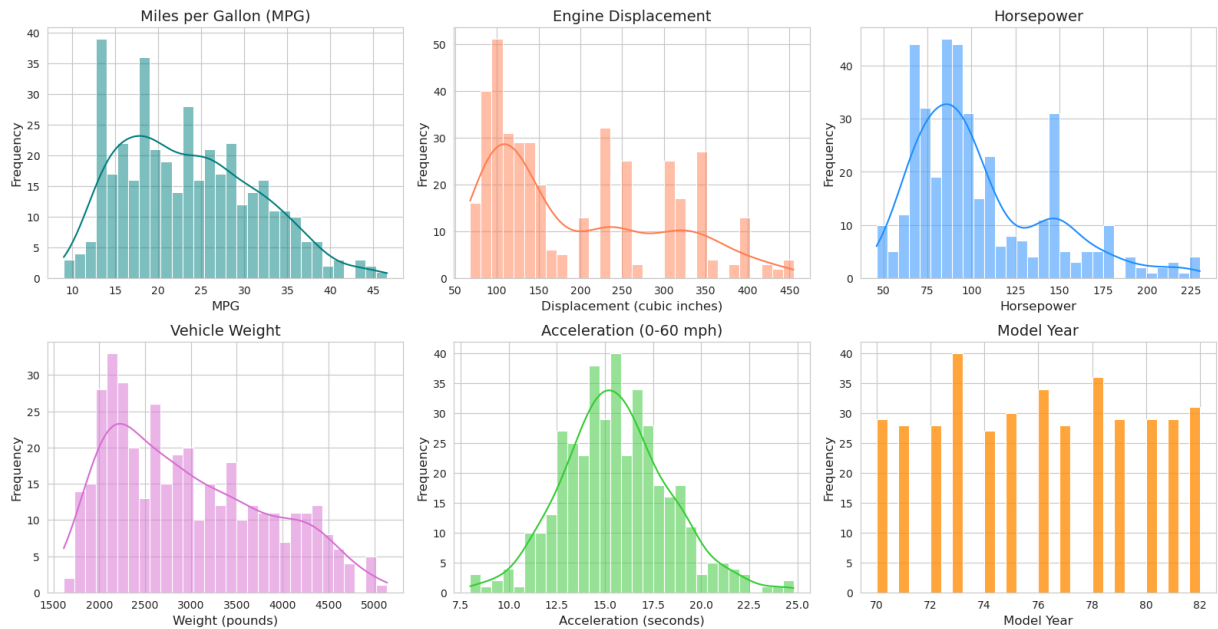
# Subplot 4 - Vehicle Weight
plt.subplot(3, 3, 4)
sns.histplot(df['weight'], bins=30, kde=True, color='orchid')
plt.title('Vehicle Weight', fontsize=14)
plt.xlabel('Weight (pounds)', fontsize=12)
plt.ylabel('Frequency', fontsize=12)

# Subplot 5 - Acceleration
plt.subplot(3, 3, 5)
sns.histplot(df['acceleration'], bins=30, kde=True, color='limegreen')
plt.title('Acceleration (0-60 mph)', fontsize=14)
plt.xlabel('Acceleration (seconds)', fontsize=12)
plt.ylabel('Frequency', fontsize=12)

# Subplot 6 - Model Year
plt.subplot(3, 3, 6)
sns.histplot(df['model year'], bins=30, kde=False, color='darkorange')
plt.title('Model Year', fontsize=14)
plt.xlabel('Model Year', fontsize=12)
plt.ylabel('Frequency', fontsize=12)

plt.tight_layout()

plt.show()
```



```
In [31]: import seaborn as sns
import matplotlib.pyplot as plt

sns.set_style("whitegrid")

plt.figure(figsize=(16, 12))

# Subplot 1 - Boxplot for MPG (Miles per Gallon)
plt.subplot(3, 3, 1)
sns.boxplot(x=df['mpg'], color='teal')
plt.title('Miles per Gallon (MPG)', fontsize=14)
plt.xlabel('MPG', fontsize=12)

# Subplot 2 - Boxplot for Displacement
plt.subplot(3, 3, 2)
sns.boxplot(x=df['displacement'], color='coral')
plt.title('Engine Displacement', fontsize=14)
plt.xlabel('Displacement (cubic inches)', fontsize=12)

# Subplot 3 - Boxplot for Horsepower
plt.subplot(3, 3, 3)
sns.boxplot(x=df['horsepower'], color='dodgerblue')
plt.title('Horsepower', fontsize=14)
plt.xlabel('Horsepower', fontsize=12)

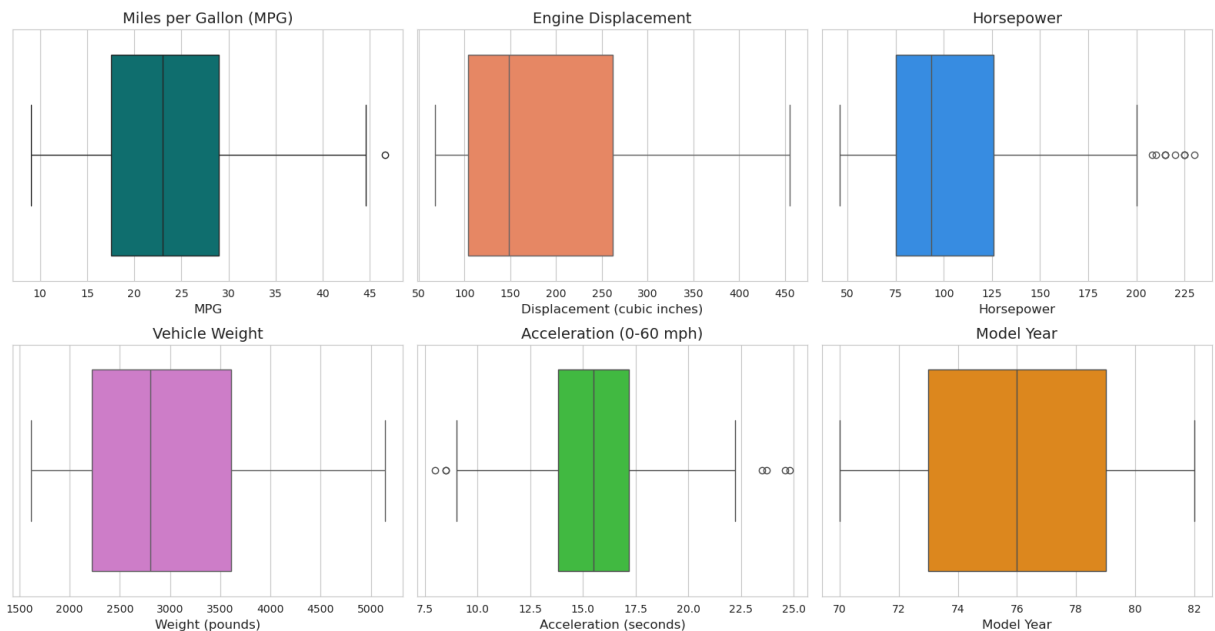
# Subplot 4 - Boxplot for Vehicle Weight
plt.subplot(3, 3, 4)
sns.boxplot(x=df['weight'], color='orchid')
plt.title('Vehicle Weight', fontsize=14)
plt.xlabel('Weight (pounds)', fontsize=12)

# Subplot 5 - Boxplot for Acceleration
plt.subplot(3, 3, 5)
sns.boxplot(x=df['acceleration'], color='limegreen')
plt.title('Acceleration (0-60 mph)', fontsize=14)
plt.xlabel('Acceleration (seconds)', fontsize=12)
```

```
# Subplot 6 - Boxplot for Model Year
plt.subplot(3, 3, 6)
sns.boxplot(x=df['model year'], color='darkorange')
plt.title('Model Year', fontsize=14)
plt.xlabel('Model Year', fontsize=12)

plt.tight_layout()

plt.show()
```



Cleaning

```
In [32]: missing_values = df.isnull().sum()
print(missing_values)
```

```
mpg          0
cylinders    0
displacement  0
horsepower    6
weight        0
acceleration  0
model year    0
origin        0
car name      0
dtype: int64
```

```
In [33]: duplicate_rows = df.duplicated().sum()
print(f"Number of duplicate rows: {duplicate_rows}")
```

Number of duplicate rows: 0

```
In [34]: from sklearn.impute import SimpleImputer

imputer = SimpleImputer(strategy='mean')
```

```
df['horsepower'] = imputer.fit_transform(df[['horsepower']])
```

```
In [35]: print(df.shape)
```

```
(398, 9)
```

```
In [36]: # IQR: Miles per Gallon (MPG), Engine Displacement, Horsepower , Vehicle Weight
# Z-Score: Acceleration (0-60 mph)
# Not Applicable: Model Year (categorical).

#first find the no of outliers in each column

from scipy.stats import zscore
from scipy.stats import skew
from scipy import stats
Iqr_columns = ['mpg', 'displacement', 'horsepower', 'weight']
Zscore_columns = ['acceleration']
def calculate_skewness(data):
    return stats.skew(data)

for column in Iqr_columns:
    skewness = calculate_skewness(df[column])
    print(f"Skewness of {column}: {skewness}")

Q1 = df[Iqr_columns].quantile(0.25)
Q3 = df[Iqr_columns].quantile(0.75)
IQR = Q3 - Q1

lower_bound = {}
upper_bound = {}

multipliers = {'mpg': 1.5, 'displacement': 1, 'horsepower': 0.9, 'weight': 0.9}
for column in Iqr_columns:
    lower_bound[column] = Q1[column] - multipliers[column] * IQR[column]
    upper_bound[column] = Q3[column] + multipliers[column] * IQR[column]

    outliers = df[(df[column] < lower_bound[column]) | (df[column] > upper_bound[column])]
    print(f"Number of outliers in {column}: {outliers}")

    df = df[(df[column] >= lower_bound[column]) & (df[column] <= upper_bound[column])]

# Calculate the Z-Score for the 'acceleration' column
df['acceleration_zscore'] = zscore(df['acceleration'])

# Find the number of outliers
outliers = df[(df['acceleration_zscore'] > 3) | (df['acceleration_zscore'] < -3)]

print(f"Number of outliers in 'acceleration': {outliers}")

# Remove the outliers
df = df[(df['acceleration_zscore'] <= 3) & (df['acceleration_zscore'] >= -3)]

df = df.drop('acceleration_zscore', axis=1)
```

Skewness of mpg: 0.45534192556309266
Skewness of displacement: 0.716930089340474
Skewness of horsepower: 1.0914191838332945
Skewness of weight: 0.5290589216608383
Number of outliers in mpg: 1
Number of outliers in displacement: 9
Number of outliers in horsepower: 23
Number of outliers in weight: 14
Number of outliers in 'acceleration': 4

```
In [37]: import seaborn as sns
import matplotlib.pyplot as plt

sns.set_style("whitegrid")

plt.figure(figsize=(16, 12))

# Subplot 1 - MPG (Miles per Gallon)
plt.subplot(3, 3, 1)
sns.histplot(df['mpg'], bins=30, kde=True, color='teal')
plt.title('Miles per Gallon (MPG)', fontsize=14)
plt.xlabel('MPG', fontsize=12)
plt.ylabel('Frequency', fontsize=12)

# Subplot 2 - Displacement
plt.subplot(3, 3, 2)
sns.histplot(df['displacement'], bins=30, kde=True, color='coral')
plt.title('Engine Displacement', fontsize=14)
plt.xlabel('Displacement (cubic inches)', fontsize=12)
plt.ylabel('Frequency', fontsize=12)

# Subplot 3 - Horsepower
plt.subplot(3, 3, 3)
sns.histplot(df['horsepower'], bins=30, kde=True, color='dodgerblue')
plt.title('Horsepower', fontsize=14)
plt.xlabel('Horsepower', fontsize=12)
plt.ylabel('Frequency', fontsize=12)

# Subplot 4 - Vehicle Weight
plt.subplot(3, 3, 4)
sns.histplot(df['weight'], bins=30, kde=True, color='orchid')
plt.title('Vehicle Weight', fontsize=14)
plt.xlabel('Weight (pounds)', fontsize=12)
plt.ylabel('Frequency', fontsize=12)

# Subplot 5 - Acceleration
plt.subplot(3, 3, 5)
sns.histplot(df['acceleration'], bins=30, kde=True, color='limegreen')
plt.title('Acceleration (0-60 mph)', fontsize=14)
plt.xlabel('Acceleration (seconds)', fontsize=12)
plt.ylabel('Frequency', fontsize=12)

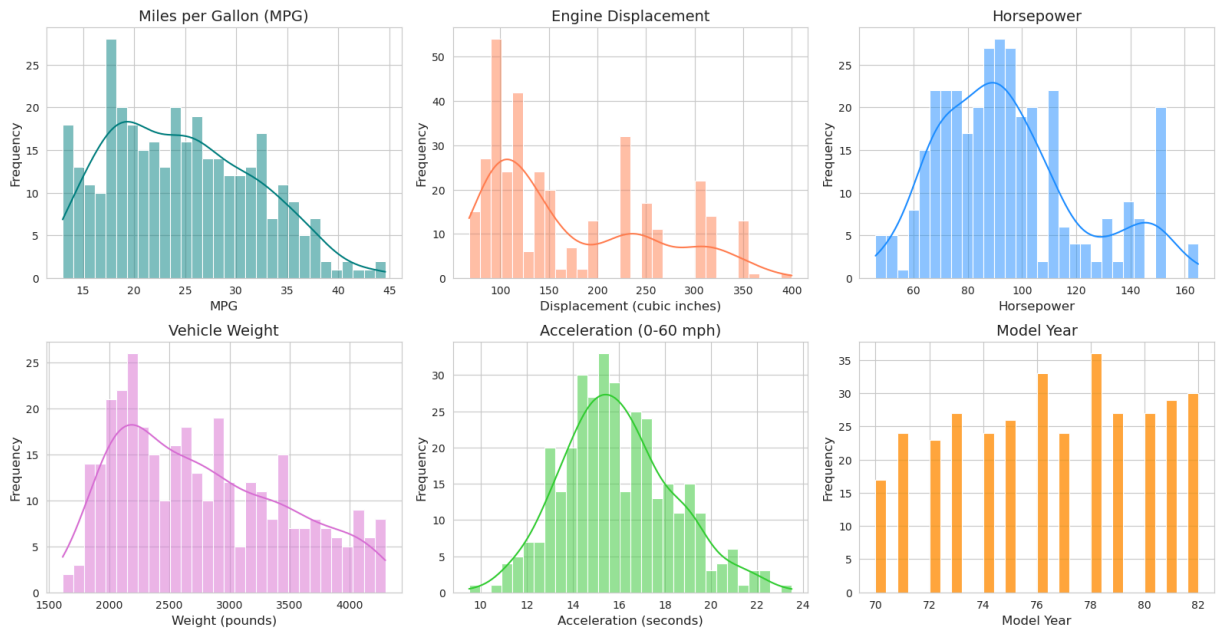
# Subplot 6 - Model Year
plt.subplot(3, 3, 6)
sns.histplot(df['model year'], bins=30, kde=False, color='darkorange')
plt.title('Model Year', fontsize=14)
```



```
plt.xlabel('Model Year', fontsize=12)
plt.ylabel('Frequency', fontsize=12)

plt.tight_layout()

plt.show()
```



```
In [38]: import seaborn as sns
import matplotlib.pyplot as plt

sns.set_style("whitegrid")

plt.figure(figsize=(16, 12))

# Subplot 1 - Boxplot for MPG (Miles per Gallon)
plt.subplot(3, 3, 1)
sns.boxplot(x=df['mpg'], color='teal')
plt.title('Miles per Gallon (MPG)', fontsize=14)
plt.xlabel('MPG', fontsize=12)

# Subplot 2 - Boxplot for Displacement
plt.subplot(3, 3, 2)
sns.boxplot(x=df['displacement'], color='coral')
plt.title('Engine Displacement', fontsize=14)
plt.xlabel('Displacement (cubic inches)', fontsize=12)

# Subplot 3 - Boxplot for Horsepower
plt.subplot(3, 3, 3)
sns.boxplot(x=df['horsepower'], color='dodgerblue')
plt.title('Horsepower', fontsize=14)
plt.xlabel('Horsepower', fontsize=12)

# Subplot 4 - Boxplot for Vehicle Weight
plt.subplot(3, 3, 4)
sns.boxplot(x=df['weight'], color='orchid')
plt.title('Vehicle Weight', fontsize=14)
```


```
plt.xlabel('Weight (pounds)', fontsize=12)

# Subplot 5 - Boxplot for Acceleration
plt.subplot(3, 3, 5)
sns.boxplot(x=df['acceleration'], color='limegreen')
plt.title('Acceleration (0-60 mph)', fontsize=14)
plt.xlabel('Acceleration (seconds)', fontsize=12)

# Subplot 6 - Boxplot for Model Year
plt.subplot(3, 3, 6)
sns.boxplot(x=df['model year'], color='darkorange')
plt.title('Model Year', fontsize=14)
plt.xlabel('Model Year', fontsize=12)

plt.tight_layout()

plt.show()
```



Encoding

```
In [39]: from sklearn.preprocessing import OneHotEncoder

import pandas as pd

encoder = OneHotEncoder(sparse_output=False)

encoded_cols = pd.DataFrame(encoder.fit_transform(df[['car name']]), columns

df = df.drop('car name', axis=1)

df = pd.concat([df, encoded_cols], axis=1)
```

```
In [43]: from sklearn.preprocessing import StandardScaler

features_to_scale = ['mpg', 'displacement', 'horsepower', 'weight', 'acceler
```

```

scaler = StandardScaler()
df_scaled = df.copy()
df_scaled[features_to_scale] = scaler.fit_transform(df[features_to_scale])

# Compare original and scaled data
for feature in features_to_scale:
    fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 4))

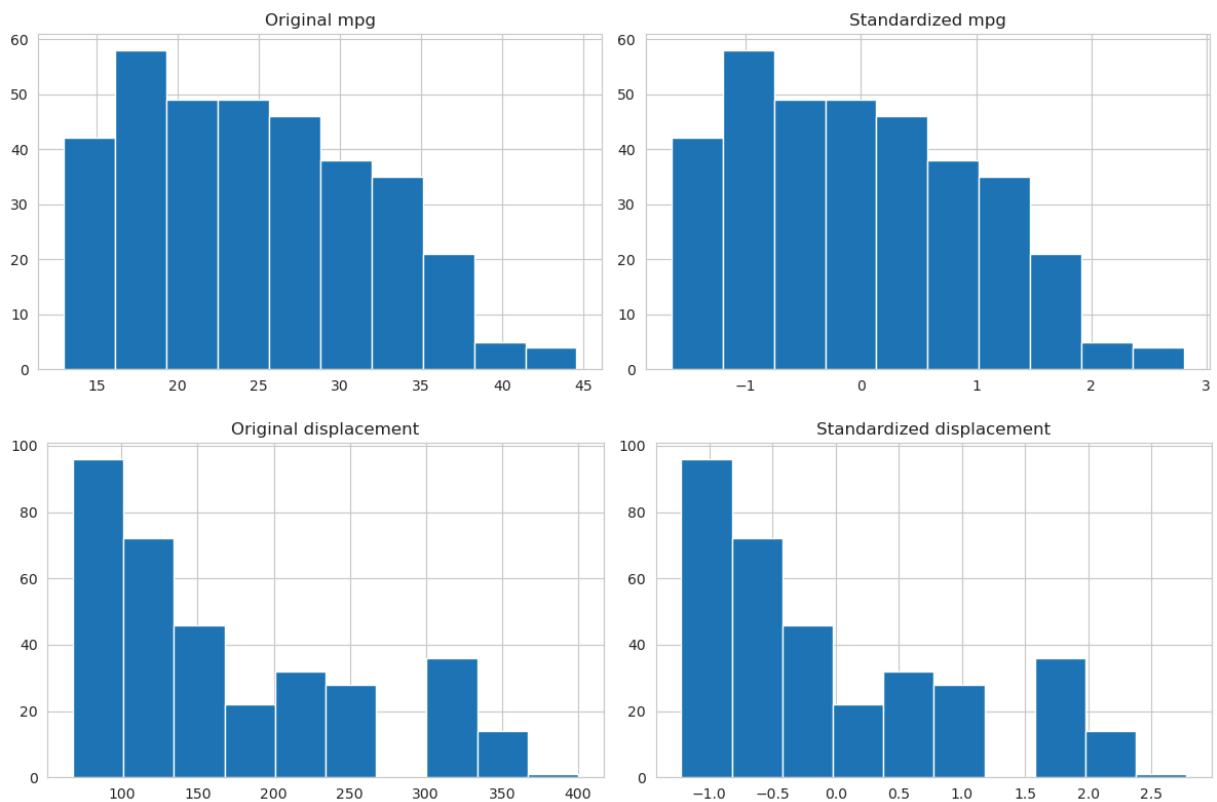
    # Original data
    df[feature].hist(ax=ax1)
    ax1.set_title(f'Original {feature}')

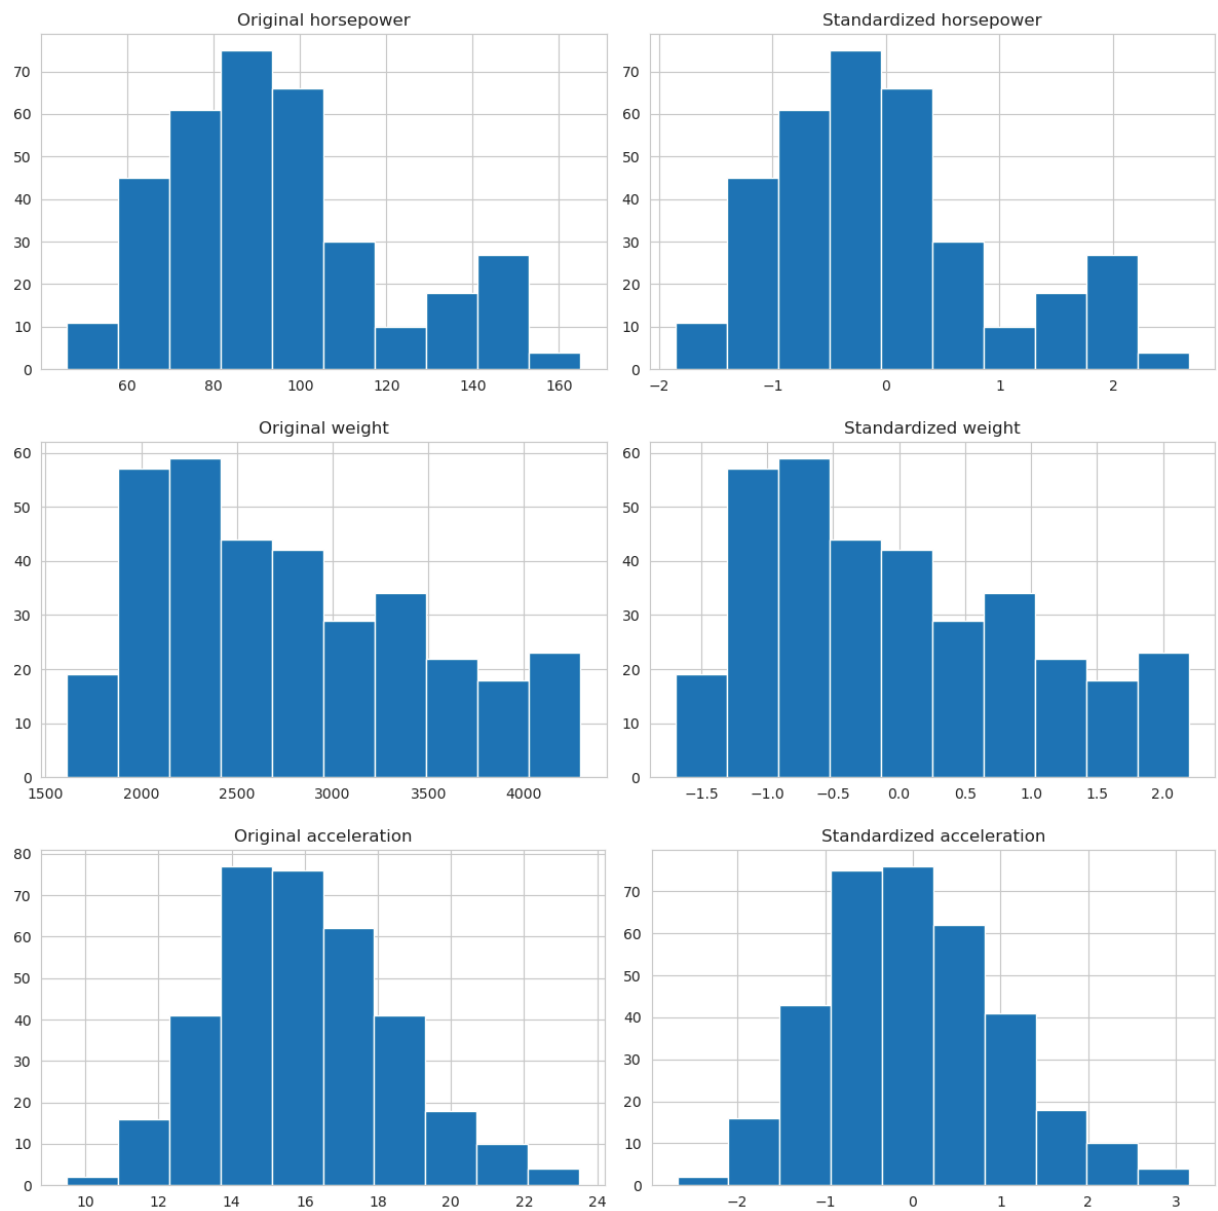
    # Scaled data
    df_scaled[feature].hist(ax=ax2)
    ax2.set_title(f'Standardized {feature}')

    plt.tight_layout()
    plt.show()

print("Original data statistics:")
print(df[features_to_scale].describe())
print("\nStandardized data statistics:")
print(df_scaled[features_to_scale].describe())

```





Original data statistics:

	mpg	displacement	horsepower	weight	acceleration
count	347.000000	347.000000	347.000000	347.000000	347.000000
mean	24.683573	169.416427	94.766041	2780.662824	15.942075
std	7.112575	83.116223	26.355453	689.889343	2.403710
min	13.000000	68.000000	46.000000	1613.000000	9.500000
25%	19.000000	98.000000	75.000000	2189.500000	14.400000
50%	24.000000	140.000000	90.000000	2665.000000	15.700000
75%	30.000000	232.000000	109.000000	3295.000000	17.400000
max	44.600000	400.000000	165.000000	4295.000000	23.500000

Standardized data statistics:

	mpg	displacement	horsepower	weight	acceleration
count	3.470000e+02	3.470000e+02	3.470000e+02	3.470000e+02	3.470000e+02
mean	-1.228604e-16	1.638139e-16	5.221568e-16	3.276278e-16	3.634621e-16
std	1.001444e+00	1.001444e+00	1.001444e+00	1.001444e+00	1.001444e+00
min	-1.645037e+00	-1.221938e+00	-1.852993e+00	-1.694980e+00	-2.683926e+00
25%	-8.002420e-01	-8.604765e-01	-7.510622e-01	-8.581325e-01	-6.424660e-01
50%	-9.624653e-02	-3.544303e-01	-1.810981e-01	-1.678963e-01	-1.008543e-01
75%	7.485480e-01	7.540519e-01	5.408563e-01	7.466123e-01	6.074071e-01
max	2.804215e+00	2.778237e+00	2.668722e+00	2.198213e+00	3.148816e+00

```
In [44]: from sklearn.preprocessing import MinMaxScaler
min_max_scaler = MinMaxScaler()
df_minmax = df.copy()
df_minmax[features_to_scale] = min_max_scaler.fit_transform(df[features_to_scale])

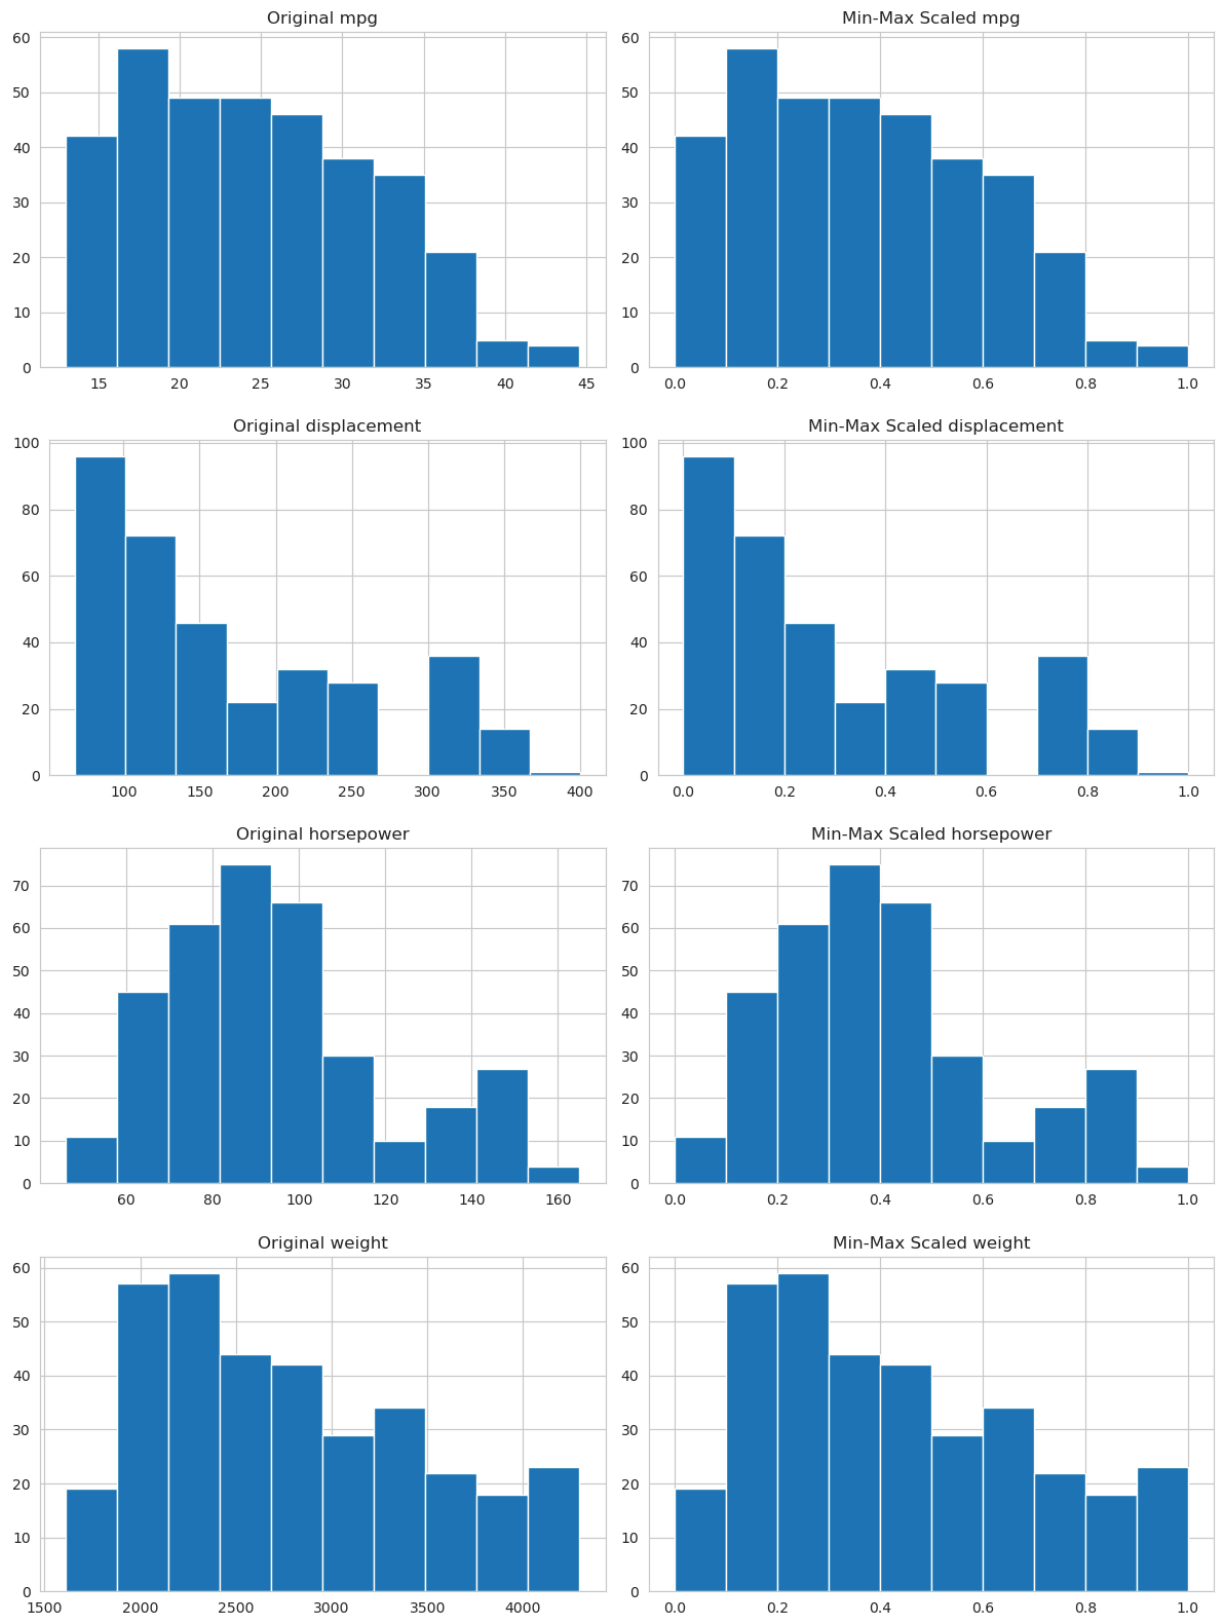
# Compare original and min-max scaled data
for feature in features_to_scale:
    fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 4))


    # Original data
    df[feature].hist(ax=ax1)
    ax1.set_title(f'Original {feature}')

    # Min-Max scaled data
    df_minmax[feature].hist(ax=ax2)
    ax2.set_title(f'Min-Max Scaled {feature}')

plt.tight_layout()
plt.show()

print("Original data statistics:")
print(df[features_to_scale].describe())
print("\nMin-Max scaled data statistics:")
print(df_minmax[features_to_scale].describe())
```





Original data statistics:

	mpg	displacement	horsepower	weight	acceleration
count	347.000000	347.000000	347.000000	347.000000	347.000000
mean	24.683573	169.416427	94.766041	2780.662824	15.942075
std	7.112575	83.116223	26.355453	689.889343	2.403710
min	13.000000	68.000000	46.000000	1613.000000	9.500000
25%	19.000000	98.000000	75.000000	2189.500000	14.400000
50%	24.000000	140.000000	90.000000	2665.000000	15.700000
75%	30.000000	232.000000	109.000000	3295.000000	17.400000
max	44.600000	400.000000	165.000000	4295.000000	23.500000

Min-Max scaled data statistics:

	mpg	displacement	horsepower	weight	acceleration
count	347.000000	347.000000	347.000000	347.000000	347.000000
mean	0.369733	0.305471	0.409799	0.435370	0.460148
std	0.225081	0.250350	0.221474	0.257229	0.171694
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.189873	0.090361	0.243697	0.214952	0.350000
50%	0.348101	0.216867	0.369748	0.392245	0.442857
75%	0.537975	0.493976	0.529412	0.627144	0.564286
max	1.000000	1.000000	1.000000	1.000000	1.000000

In []: