

Name : Tazmeen Afroz

Roll No : 22P-9252

Section : BAI-5A

Lab Task 07

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

```
In [3]: def compute_cost(x, y, w, b):
        """
        Computes the cost function for linear regression.
        Args:
        x (ndarray (m,)): Data, m examples
        y (ndarray (m,)): target values
        w,b (scalar) : model parameters
        Returns
        total_cost (float): The cost of using w,b as the parameters for linear r
        to fit the data points in x and y
        """
        m = x.shape[0]
        cost_sum = 0
        for i in range(m):
            f_wb = w * x[i] + b
            cost = (f_wb - y[i]) ** 2
            cost_sum = cost_sum + cost
        total_cost = (1 / (2 * m)) * cost_sum
        return total_cost
```

```
In [4]: # Load the data
df = pd.read_csv('housing.csv')
print(df.head())
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-122.23	37.88	41.0	880.0	129.0	
1	-122.22	37.86	21.0	7099.0	1106.0	
2	-122.24	37.85	52.0	1467.0	190.0	
3	-122.25	37.85	52.0	1274.0	235.0	
4	-122.25	37.85	52.0	1627.0	280.0	

	population	households	median_income	median_house_value	ocean_proximity
0	322.0	126.0	8.3252	452600.0	NEAR BAY
1	2401.0	1138.0	8.3014	358500.0	NEAR BAY
2	496.0	177.0	7.2574	352100.0	NEAR BAY
3	558.0	219.0	5.6431	341300.0	NEAR BAY
4	565.0	259.0	3.8462	342200.0	NEAR BAY

```
In [5]: X = df['housing_median_age']
        y = df['median_income']

        print(X.head())
        print(y.head())
```

```
0    41.0
1    21.0
2    52.0
3    52.0
4    52.0
Name: housing_median_age, dtype: float64
0    8.3252
1    8.3014
2    7.2574
3    5.6431
4    3.8462
Name: median_income, dtype: float64
```

```
In [6]: X.info()
```

```
<class 'pandas.core.series.Series'>
RangeIndex: 20640 entries, 0 to 20639
Series name: housing_median_age
Non-Null Count  Dtype
-----
20640 non-null  float64
dtypes: float64(1)
memory usage: 161.4 KB
```

```
In [7]: # null values check
```

```
print(X.isnull().sum())
print(y.isnull().sum())
```

```
0
0
```

```
In [8]: # scaling
```

```
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
```


```
X = scaler.fit_transform(X.values.reshape(-1, 1))
y = scaler.fit_transform(y.values.reshape(-1, 1))

print(X[:5])
```

```
[[ 0.98214266]
 [-0.60701891]
 [ 1.85618152]
 [ 1.85618152]
 [ 1.85618152]]
```

```
In [9]: # plotting the data
import matplotlib.pyplot as plt

plt.scatter(X, y)
plt.xlabel('housing_median_age')
plt.ylabel('median_income')
plt.show()
```



```
In [10]: X_train, X_test, y_train, y_test = train_test_split(X_scaled, y_scaled, test_size=0.2, random_state=42)

# Linear regression model
regr = LinearRegression()
regr.fit(X_train, y_train)

# Make predictions
y_pred = regr.predict(X_test)
```

```
In [11]: mse = mean_squared_error(y_test, y_pred)
```

```
In [12]: print(f"Mean Squared Error: {mse}")
```

Mean Squared Error: 0.9538407229513474

```
In [13]: # Calculate custom cost
w = regr.coef_[0]
b = regr.intercept_
custom_cost = compute_cost(X_test.flatten(), y_test, w, b)
print(f"Custom Cost: {custom_cost}")
```

Custom Cost: 0.47692036147567546

```
In [14]: # Plot the regression line
plt.figure(figsize=(10, 6))
plt.scatter(X_test, y_test, color='blue', label='Actual')
plt.plot(X_test, y_pred, color='red', label='Predicted')
plt.xlabel('Scaled Housing Median Age')
plt.ylabel('Scaled Median Income')
plt.title('Linear Regression: Housing Median Age vs Median Income')
plt.legend()
plt.show()
```



```
In [ ]:
```